

CEREBRUM

Community-Structured Graph Attention for Knowledge Graph Reasoning

MARKET ANALYSIS REPORT

Confidential — Prepared for Strategic and Investment Review

Developer: Bryan Alexander Buchorn (© 2026)

Report Date: March 2026

Classification: Confidential / Pre-Commercial

1. Executive Summary

Core Thesis

The dominant trend in enterprise AI — scaling large language models — produces systems that are powerful but structurally unverifiable. In every regulated, high-stakes domain, unverifiability is a disqualifying flaw. CEREBRUM represents a structural alternative: reasoning that is competitive in accuracy, verifiable by design, and deployable immediately without training or GPUs.

CEREBRUM is a novel AI reasoning framework that performs multi-hop traversal over knowledge graphs (KGs) using Transformer-analogous structural principles — without large language models, without training data, and without GPU infrastructure. Developed by Bryan Alexander Buchorn, the system introduces two architectural primitives: the Dual-Signal Community Fusion (DSCF) algorithm and Community-Structured Attention (CSA) mechanism.

In benchmark testing against the MetaQA dataset (43,234 entities / 124,680 edges), CEREBRUM achieves results that are competitive with — and in some configurations superior to — trained systems on 1-hop and 2-hop reasoning tasks, while requiring zero training cost. At 3-hop depth, a performance gap exists relative to retrieval-augmented systems; this is a known and quantifiable risk discussed in Section 8.

The commercial opportunity is large and specific. The global knowledge graph market is projected to exceed \$10.7 billion by 2030. Within that, regulated industries — healthcare, finance, legal, and defense — represent a serviceable addressable market of approximately \$4 to \$6 billion where verifiable, auditable AI reasoning is not a preference but a compliance requirement. CEREBRUM occupies the only position on the competitive landscape that satisfies all four criteria simultaneously: high accuracy, zero training, full verifiability, and hardware-agnostic deployment.

Revenue is structured as a dual license: AGPL-3.0 for community and academic use, and a commercial license for enterprise and SaaS deployment. This model creates a developer community flywheel while protecting commercial upside. Under base-case projections, CEREBRUM reaches \$3.2 million in annual recurring revenue by Year 3, with an optimistic case of \$12 million ARR if enterprise healthcare or financial services adoption accelerates.

Summary Metrics

Metric	Assessment
Market Opportunity (TAM)	\$10.7B KG market by 2030 (28% CAGR); \$100B+ enterprise AI infrastructure market
Serviceable Market (SAM)	\$4–6B (regulated industries requiring verifiable AI: healthcare, finance, legal, defense)
Year 3 ARR — Base Case	\$3.2M (conservative: \$0.8M; optimistic: \$12M)
Competitive Position	Only system achieving high accuracy + zero training + full audit trail + CPU-compatible
Key Regulatory Tailwinds	EU AI Act (Art. 13), FDA 21 CFR Part 11, SOX, HIPAA, FINRA — all favor verifiable AI
Primary Risk	3-hop accuracy gap vs retrieval-augmented systems; addressable via hybrid pipeline
Go-to-Market	OSS community flywheel + ArXiv/academic credibility + targeted enterprise vertical sales
Licensing	AGPL-3.0 (free) + Commercial license (enterprise/SaaS)

2. Total Addressable Market (TAM / SAM / SOM)

2.1 Total Addressable Market (TAM)

CEREBRUM sits at the intersection of three distinct but converging markets: knowledge graph platforms, enterprise AI infrastructure, and explainable/auditable AI tooling.

Knowledge Graph Market

The global knowledge graph market was valued at approximately \$1.1 billion in 2023 and is projected to reach \$10.7 billion by 2030, representing a compound annual growth rate (CAGR) of 28.4% (Source: Grand View Research, MarketsandMarkets — flagged as external estimates). Growth is driven by enterprise demand for structured data integration, semantic search, and multi-source data federation.

Enterprise AI Infrastructure Market

Enterprise AI infrastructure — including inference platforms, orchestration layers, and reasoning tools — was valued at approximately \$50 billion in 2024 and is expected to exceed \$220 billion by 2030. Within this, the segment focused on structured reasoning and knowledge-grounded AI (as opposed to generative text) represents roughly 8–12% of total spend, or \$4–6 billion today growing to \$25–30 billion by 2030. This estimate reflects the increasing separation of use cases between generative AI (content creation, summarization) and structured AI (compliance, clinical decision support, financial reasoning). (Assumption flagged: segment breakdown is analyst estimate.)

Explainable / Auditable AI Market

Explainable AI (XAI) is an emerging sub-segment driven primarily by regulatory pressure. The global XAI market was valued at \$6.5 billion in 2023 and is projected to reach \$24.7 billion by 2030 (CAGR ~21%). CEREBRUM's audit-trail architecture positions it squarely within this segment as a native solution rather than a post-hoc explanation layer applied to a black-box model.

Market	2024 Size	2030 Projection	CAGR
Knowledge Graph Platforms	~\$1.5B	~\$10.7B	28.4%
Enterprise AI Infrastructure	~\$50B	~\$220B	~28%
Structured Reasoning Segment	~\$5B	~\$27B	~28%
Explainable AI	~\$6.5B	~\$24.7B	~21%
Combined TAM (overlap-adjusted)	~\$15–20B	~\$60–80B	~25–28%

2.2 Serviceable Addressable Market (SAM)

CEREBRUM's SAM is defined by domains where verifiable, non-hallucinating AI reasoning is a hard requirement — not a differentiator but a baseline for deployment. These are regulated

environments where an incorrect AI inference carries legal, financial, or human health consequences.

- **Healthcare / Clinical Decision Support** — AI-assisted diagnosis, drug interaction checking, clinical pathway navigation. Regulatory driver: FDA 21 CFR Part 11, ONC interoperability rules. Market: \$45B total health AI market by 2026; relevant structured reasoning subset approximately \$3–5B.
- **Financial Services / Risk & Compliance** — Regulatory reporting, anti-money laundering (AML) graph analysis, credit risk reasoning. Regulatory driver: SOX, FINRA, Basel III. Market: \$30B+ financial AI market; structured risk reasoning subset approximately \$4–6B.
- **Legal / Contract Intelligence** — Contract clause extraction, obligation mapping, regulatory compliance checking. Market: \$2.5B legal AI market growing at 32% CAGR.
- **Defense & Intelligence** — Threat graph analysis, OSINT relationship mapping, supply chain risk. Procurement driven by explainability requirements in AI acquisition frameworks (DoD AI Strategy). Market: \$9B+ defense AI market with structured analytics as a core procurement category.
- **Drug Discovery / Biomedical Research** — Gene-disease relationship traversal, protein interaction networks, adverse event detection. Market: \$4B+ in AI for drug discovery.

Aggregate SAM estimate: \$12–18 billion in 2024, growing to \$50–65 billion by 2030. The binding constraint on SAM is not demand — it is supply of verifiable AI systems. That is the gap CEREBRUM fills.

2.3 Serviceable Obtainable Market (SOM)

SOM reflects what a two-person founding team with a dual-license model can realistically capture in Years 1 through 3. The model assumes: (a) no institutional sales team in Year 1, (b) organic developer adoption driving enterprise pipeline, (c) first enterprise deals closing in Year 2.

Year	Conservative	Base Case	Optimistic
Year 1	\$80K ARR	\$180K ARR	\$400K ARR
Year 2	\$420K ARR	\$900K ARR	\$3.5M ARR
Year 3	\$820K ARR	\$3.2M ARR	\$12M ARR

Key assumptions: conservative scenario assumes 2–3 small enterprise deals by Year 3; base case assumes one anchor vertical (healthcare or finance) with 5–8 commercial licenses; optimistic case assumes a cloud marketplace listing driving inbound volume and one OEM partnership.

3. Competitive Landscape

3.1 Direct Competitors

Graph Database Vendors

Neo4j, TigerGraph, and Amazon Neptune are graph storage and query systems, not reasoning engines. They answer structured Cypher or Gremlin queries but do not perform multi-hop inference or provide reasoning paths as a first-class output. They are platforms CEREBRUM can run on top of (via the THALAMUS ingestion adapter) rather than competitors.

- Neo4j: \$300M+ ARR, 4,000+ customers. Strength: mature graph DB with broad enterprise adoption. Gap: no native reasoning layer, no verifiable inference path, no training-free KG QA.
- TigerGraph: Focused on analytics graphs (fraud, supply chain). Offers graph ML capabilities but requires model training. No architectural equivalent to CSA or DSCF.
- Amazon Neptune: Managed graph DB in AWS ecosystem. Strong integration story but passive storage layer. Amazon's own ML on Neptune requires SageMaker training pipelines.

Trained KG Reasoning Systems

EmbedKGQA, MINERVA, GraftNet, and PullNet are the primary academic benchmarks CEREBRUM is measured against. All require significant training investment and, in the case of GraftNet and PullNet, external text retrieval to compensate for graph coverage gaps.

System	1-hop H@10	2-hop H@10	3-hop H@10	Training	External Text
CEREBRUM FULL	96.9%	72.9%	38.3%	None	No
EmbedKGQA	95.0%	68.6%	49.6%	Required	No
MINERVA	95.3%	78.2%	45.6%	Required	No
GraftNet	93.9%	77.3%	72.1%	Required	Yes
PullNet	97.0%	89.6%	81.8%	Required	Yes

Note: GraftNet and PullNet's strong 3-hop performance depends on external text retrieval to bridge graph gaps. These systems are not pure structural reasoners. In closed-graph environments (healthcare records, classified data, proprietary financial graphs), text retrieval pipelines introduce data exposure risk — an environment where CEREBRUM's closed-graph architecture is a security advantage.

RAG Platforms

LlamaIndex and LangChain are orchestration frameworks for retrieval-augmented generation — they route queries to LLMs with retrieved context. They are architecturally dependent on an underlying LLM and inherit all of LLMs' hallucination risk. They are not graph reasoning systems. The comparison is primarily relevant in enterprise evaluation contexts where CEREBRUM competes against a RAG pipeline for knowledge retrieval tasks.

3.2 Indirect Competitors — LLM Providers

OpenAI, Anthropic, and Google represent the most significant indirect competitive threat. Each has built or is actively developing structured reasoning capabilities layered on top of their base LLMs (e.g., OpenAI's tool-use APIs, Google's NotebookLM structured retrieval, Anthropic's Projects). However, these remain probabilistic systems — they do not produce verifiable traversal paths, cannot guarantee hallucination-free outputs, and require ongoing inference compute and API costs. In regulated environments, LLM-based systems face the additional challenge that their outputs cannot be traced to a specific evidence path, which undermines audit trail requirements.

3.3 Competitive Positioning — 2x2 Map

The following matrix maps the competitive landscape on two dimensions: (a) output verifiability (can the system trace its answer to specific evidence?) and (b) training cost (does deployment require data collection, labeling, and model training?).

	Low Training Cost	High Training Cost	
High Verifiability	CEREBRUM ★ (whitespace)	EmbedKGQA, MINERVA	
Low Verifiability	LlamaIndex, LangChain (RAG)	OpenAI, Anthropic, Google (LLMs)	

Whitespace Finding

No commercially available system occupies the High Verifiability + Low Training Cost quadrant at production scale. CEREBRUM is the only system in this position with documented benchmark performance competitive with trained systems. This is not positioning — it is a structural property of the architecture.

4. Industry Verticals — Impact Analysis

4.1 Healthcare / Clinical Decision Support

- **Pain Point:** Clinical AI systems operating in diagnostic or treatment recommendation roles must produce explainable outputs. Under FDA guidance and ONC interoperability rules, a system that returns a recommendation without an auditable reasoning path cannot be deployed in a regulated clinical workflow. LLM-based clinical decision tools are blocked from many use cases precisely because they cannot trace outputs to specific evidence.
- **Regulatory Driver:** FDA 21 CFR Part 11 requires electronic records and audit trails for software used in clinical decisions. The EU AI Act classifies clinical decision support as high-risk AI (Annex III), requiring logging, transparency, and human oversight. CEREBRUM's edge-traced traversal paths satisfy both requirements natively.
- **Buyer Persona:** VP of Clinical Informatics, Chief Medical Information Officer (CMIO), or Director of Regulatory Affairs at a hospital system, health IT vendor, or EHR platform provider (Epic, Cerner/Oracle Health, athenahealth).
- **Estimated Deal Size:** \$80K–\$250K/year for hospital system deployment; \$300K–\$1M+ for EHR platform OEM integration. Integration with an EHR vendor's platform represents the highest-leverage single deal in this vertical.

4.2 Financial Services / Risk & Compliance

- **Pain Point:** Regulatory reporting, anti-money laundering (AML) entity graph analysis, and credit risk assessment all require traceable decision paths. A compliance officer cannot submit a FINRA examination response that references an AI recommendation without an audit trail. Currently, most financial AI systems either use rule-based engines (auditable but brittle) or LLMs (flexible but unverifiable).
- **Regulatory Driver:** SOX Section 404 requires documented and auditable internal controls. FINRA Rule 3110 requires supervisory systems for AI-assisted recommendations. Basel III and forthcoming Basel IV frameworks incorporate AI model risk management requirements. All favor deterministic, traceable reasoning.
- **Buyer Persona:** Chief Risk Officer, Head of Compliance Technology, or RegTech procurement team at a bank, hedge fund, or insurance company.
- **Estimated Deal Size:** \$120K–\$500K/year for compliance analytics deployment; \$1M+ for core risk model integration at a tier-1 bank.

4.3 Legal / Contract Intelligence

- **Pain Point:** Legal AI tools (e.g., Harvey, Ironclad, Kira) extract and reason over contract clauses but rely on LLMs for the reasoning layer. Hallucinated clause interpretations in M&A due diligence or contract negotiations carry direct financial and legal liability. Firms need systems that can trace a conclusion ("this clause conflicts with Regulation X") to a specific graph path.
- **Regulatory Driver:** Attorney professional conduct rules (ABA Model Rule 1.1) require competence, including understanding of AI tool limitations. Malpractice exposure drives demand for verifiable AI outputs. The EU AI Act's high-risk classification may extend to legal AI systems used in regulatory compliance contexts.
- **Buyer Persona:** Director of Legal Technology, General Counsel, or Legal Operations Manager at an Am Law 100 firm or Fortune 500 legal department.

- **Estimated Deal Size:** \$25K–\$80K/year for boutique or mid-market firm; \$200K–\$600K/year for enterprise legal department integration.

4.4 Defense & Intelligence

- **Pain Point:** Threat graph analysis, OSINT relationship mapping, and supply chain risk assessment all involve multi-hop inference over structured data. Defense procurement increasingly requires AI systems to operate in air-gapped, classified environments — which eliminates cloud-based LLMs and most RAG pipelines. CEREBRUM's CPU-only, offline deployment capability is a direct answer to this constraint.
- **Regulatory Driver:** DoD AI Strategy (2023) and DoD Directive 3000.09 require human-machine teaming and system transparency for AI-enabled decision-making in lethal and high-consequence applications. Federal AI procurement frameworks increasingly mandate explainability. CEREBRUM's architecture satisfies both.
- **Buyer Persona:** Program Manager at a defense contractor (Palantir, Leidos, SAIC, Booz Allen), or direct government customer via SBIR/STTR mechanisms.
- **Estimated Deal Size:** \$500K–\$3M per program; SBIR Phase II contracts typically \$1–2M. Direct DoD contracts at scale can reach \$10M+.

4.5 Drug Discovery / Biomedical Research

- **Pain Point:** Biomedical knowledge graphs (e.g., PrimeKG, Hetionet, STRING) contain millions of gene-disease, protein-interaction, and drug-target relationships. Identifying multi-hop paths (e.g., Drug A → inhibits → Protein B → expressed in → Disease C) requires reasoning systems that can traverse large graphs without retraining for each new research question. LLMs hallucinate biological relationships; CEREBRUM is bounded by what the graph actually contains.
- **Regulatory Driver:** FDA's emerging guidance on AI/ML-based software as a medical device (SaMD) requires traceability of AI-assisted drug discovery decisions. CEREBRUM's audit trails support regulatory filing documentation.
- **Buyer Persona:** Director of Computational Biology, Bioinformatics Lead, or Head of AI at a pharmaceutical company or biotech startup.
- **Estimated Deal Size:** \$60K–\$200K/year for research team access; \$500K–\$2M for integration into a drug discovery platform (e.g., Schrödinger, Recursion Pharmaceuticals).

4.6 Enterprise Knowledge Management

- **Pain Point:** Large organizations accumulate institutional knowledge in fragmented systems — SharePoint, Confluence, Salesforce, Jira, internal wikis. Current AI approaches (RAG chatbots) retrieve documents but cannot reason across multi-source relationships. CEREBRUM enables a graph of organizational knowledge where multi-hop queries ("Which customers use Product X and have open support tickets related to Feature Y?") return verifiable paths rather than probabilistic answers.

- **Regulatory Driver:** Less directly regulated than the above verticals, but GDPR, CCPA, and internal data governance requirements create demand for AI systems that can demonstrate what data was used in a decision — which CEREBRUM provides natively.
- **Buyer Persona:** Chief Knowledge Officer, VP of Enterprise Architecture, or Head of AI Strategy at a Fortune 500 enterprise.
- **Estimated Deal Size:** \$30K–\$120K/year for departmental deployment; \$500K+ for enterprise-wide knowledge graph platform integration.

Vertical Summary Table

Vertical	Deal Size	Regulatory Driver	Buyer	Priority
Healthcare / CDS	\$80K–\$1M+	FDA 21 CFR Part 11, EU AI Act	CMIO, VP Informatics	★★★★★
Financial / Compliance	\$120K–\$1M+	SOX, FINRA, Basel III	CRO, Compliance Tech	★★★★★
Legal / Contract	\$25K–\$600K	Malpractice, EU AI Act	Legal Ops, GC	★★★★☆
Defense & Intel	\$500K–\$3M+	DoD AI Strategy, SBIR	Program Manager	★★★★★
Drug Discovery	\$60K–\$2M	FDA SaMD guidance	Comp. Biology Director	★★★★★
Enterprise KM	\$30K–\$500K+	GDPR, internal governance	CKO, VP Architecture	★★★★☆

5. Regulatory Tailwinds

CEREBRUM's architecture was not designed to satisfy regulations — it was designed to reason correctly. The fact that it satisfies an expanding set of regulatory requirements is a structural consequence of how it works, not a compliance feature bolted on afterward. This distinction is commercially significant: competing systems must build explainability layers on top of probabilistic models; CEREBRUM's audit trail is the model.

5.1 EU Artificial Intelligence Act

The EU AI Act, which entered into force in August 2024 with enforcement phases through 2027, establishes the most comprehensive regulatory framework for AI globally. Key provisions relevant to CEREBRUM:

- **Article 13 — Transparency:** High-risk AI systems must be designed to ensure that their operation is sufficiently transparent to enable users to interpret outputs and use them appropriately. CEREBRUM's traversal paths are inherently transparent — every output is a labeled sequence of edges through the graph, each edge corresponding to a verifiable relationship.
- **Article 14 — Human Oversight:** High-risk systems must enable human supervisors to understand system behavior and detect errors. CEREBRUM's audit trail directly enables this requirement without additional tooling.
- **Article 17 — Quality Management:** Providers must establish documentation systems covering training data (if any), system design, and testing. CEREBRUM's training-free architecture eliminates an entire documentation category (training data provenance), simplifying Article 17 compliance.
- **Annex III — High-Risk Classifications:** Includes medical devices, employment screening, credit evaluation, and law enforcement AI. CEREBRUM is positioned precisely in these verticals. Systems in these categories face the highest compliance burden — and CEREBRUM satisfies it architecturally.

5.2 US Executive Orders and Federal AI Policy

The October 2023 Executive Order on AI Safety and the subsequent OMB guidance (March 2024) require federal agencies to conduct AI impact assessments, maintain documentation of AI-assisted decisions, and implement human oversight for consequential AI use cases. CEREBRUM's audit trail is compliant with these requirements. Defense and civilian agency procurement increasingly reflects these mandates.

5.3 FDA 21 CFR Part 11

21 CFR Part 11 establishes requirements for electronic records and electronic signatures used in FDA-regulated industries. Specifically, it requires that computer systems used in clinical or regulated contexts maintain complete, accurate, and accessible audit trails. A clinical decision support system built on CEREBRUM produces a native audit trail for every inference — the traversal path is the required record. Competing LLM-based systems cannot satisfy this requirement without additional (and architecturally awkward) logging infrastructure.

5.4 SOX, HIPAA, FINRA

- **SOX (Sarbanes-Oxley):** Section 404 requires management to assess and document internal controls over financial reporting. AI-assisted financial analysis tools must demonstrate their reasoning is auditable. CEREBRUM's paths satisfy this directly.
- **HIPAA:** The HIPAA Security Rule requires covered entities to implement access controls and audit controls on electronic protected health information (ePHI) systems. CEREBRUM's graph operates on structured data; its traversal logs provide a native audit record of what patient data was accessed in a clinical inference.

- **FINRA:** FINRA Regulatory Notice 22-09 and related guidance require firms using AI in customer-facing or supervisory contexts to document and test algorithmic systems. CEREBRUM's verifiable paths constitute the documentation the guidance requires.

Regulatory Compliance Summary

Regulation	Requirement	CEREBRUM Compliance Mechanism
EU AI Act Art. 13	Transparency of high-risk AI outputs	Edge-labeled traversal path is the output
EU AI Act Art. 14	Human oversight of AI decisions	Traversal paths are human-readable and reviewable
FDA 21 CFR Part 11	Electronic audit trails for regulated SW	Every inference generates a native audit log
HIPAA Security Rule	Access and audit controls on ePHI systems	Graph traversal logs document data access per query
SOX Section 404	Auditable internal controls for financial reporting	Reasoning paths are auditable financial decision records
FINRA Rule 3110	Supervisory documentation of AI-assisted decisions	Verifiable paths satisfy supervisory review requirement
DoD AI Directive 3000.09	Transparency in AI-enabled decision-making	Air-gap compatible; CPU-only; full path audit

6. Go-To-Market Strategy Assessment

6.1 Open-Source Community Flywheel

The AGPL-3.0 license serves a specific strategic function: it creates a high-friction commercial conversion mechanism. Any organization that wants to deploy CEREBRUM in a proprietary product or internal system without open-sourcing that product must acquire a commercial license. This is the same dual-license structure used successfully by MongoDB, Elasticsearch (pre-SSPL), MariaDB, and Grafana Labs.

The flywheel works as follows: (1) developers discover CEREBRUM on GitHub, PyPI, or Hugging Face; (2) they integrate it into research and proof-of-concept projects under AGPL; (3) when their organization wants to deploy commercially, legal review triggers a commercial license inquiry; (4) enterprise sales proceeds from an already-validated technical champion inside the organization. This removes cold-outreach from the early go-to-market entirely.

6.2 Academic and Research Credibility Pipeline

- **ArXiv Preprint:** Publishing a technical paper on CEREBRUM (DSCF algorithm, CSA mechanism, MetaQA benchmark results) establishes a citable precedent before commercial launch. ArXiv preprints in ML receive significant traffic and are indexed by semantic search tools used by enterprise AI researchers.
- **Papers With Code Leaderboard:** Submitting CEREBRUM results to the Papers With Code MetaQA leaderboard creates a persistent, publicly visible credibility signal. Researchers evaluating KG reasoning systems routinely check leaderboard rankings. A top-5 position on 1-hop and 2-hop MetaQA, achieved without training, is a genuinely novel result.
- **Conference Targeting:** Primary targets: ISWC (International Semantic Web Conference), ACL (knowledge graphs track), IEEE BigData, NeurIPS (structured reasoning workshops). Secondary: HIMSS (healthcare IT), Money20/20 (financial services), CDO Summit (enterprise data).

6.3 Partnership Channels

- **Cloud Marketplaces:** AWS Marketplace, Azure Marketplace, and GCP Marketplace listings provide immediate access to enterprise procurement teams with existing cloud spend commitments. Marketplace listings also trigger co-sell opportunities with cloud provider partner teams, which function as embedded sales support.
- **System Integrators:** Deloitte, Accenture, and IBM Global Services are the primary implementation partners for enterprise AI projects in healthcare, finance, and government. A technology partnership with one tier-1 SI accelerates enterprise sales by embedding CEREBRUM into existing client engagement workflows.
- **Healthcare IT Ecosystem:** Partnerships with EHR platforms (Epic App Orchard, Oracle Health Marketplace, athenahealth Marketplace) provide direct access to hospital system buyers without cold outreach. Epic alone has 28% of the US acute care hospital market.
- **Graph Database Vendors:** Neo4j's partner ecosystem includes hundreds of ISVs and consultancies. A CEREBRUM-on-Neo4j integration, published to the Neo4j Marketplace, creates a distribution channel through an established enterprise customer base.

6.4 Conference and Publication Roadmap

Venue	Timeline	Purpose
ArXiv preprint submission	Month 1–2	Establishes citable precedent; researcher discovery
Papers With Code submission	Month 2	Leaderboard credibility; inbound researcher interest
GitHub public release	Month 2–3	Developer adoption begins; AGPL flywheel activates
ISWC submission	Month 3–4	Primary academic venue for KG research

Venue	Timeline	Purpose
NeurIPS workshop submission	Month 6	ML community exposure; structured reasoning track
AWS Marketplace listing	Month 6–9	Enterprise procurement access; co-sell activation
HIMSS / Money20/20	Year 2	Enterprise buyer exposure in healthcare and finance

7. Revenue Model & Projections

7.1 License Tiers

Tier	Price	Target	Notes
Community (AGPL)	Free	Researchers, OSS projects	AGPL requires source disclosure on deployment; drives commercial conversion
Startup / Developer	\$1,000/month	Startups, small teams	Proprietary deployment without AGPL disclosure; includes support
Commercial	\$5,000/month	Mid-market enterprises	Multi-graph, multi-user; SLA support; cloud or on-prem
Enterprise	\$100,000/year	Large enterprises	Unlimited graphs; dedicated support; custom SLA; audit tooling
OEM / Embedded	Custom (rev share or flat)	Platform vendors, ISVs	CEREBRUM embedded in third-party product; highest-margin tier

7.2 Revenue Projections — Three Scenarios

All projections assume a two-person team (technical founder + part-time business development) with no institutional sales force in Year 1. Dollar figures are ARR (Annual Recurring Revenue).

Conservative Scenario

Limited enterprise traction in Year 1-2. Community adoption is moderate. First enterprise deal closes in late Year 2.

Tier	Year 1	Year 2	Year 3
Startup (\$1K/mo)	2 customers = \$24K	5 customers = \$60K	10 customers = \$120K

Tier	Year 1	Year 2	Year 3
Commercial (\$5K/mo)	0 = \$0	2 customers = \$120K	4 customers = \$240K
Enterprise (\$100K/yr)	0 = \$0	1 customer = \$100K	4 customers = \$400K
OEM	0	0	1 deal = \$60K
Total ARR	\$24K	\$280K	\$820K

Base Case Scenario

Moderate developer adoption in Year 1. One anchor enterprise vertical (healthcare or finance) closes in Year 2. ArXiv paper drives inbound interest.

Tier	Year 1	Year 2	Year 3
Startup (\$1K/mo)	6 customers = \$72K	15 customers = \$180K	25 customers = \$300K
Commercial (\$5K/mo)	1 customer = \$60K	4 customers = \$240K	8 customers = \$480K
Enterprise (\$100K/yr)	0	2 customers = \$200K	8 customers = \$800K
OEM	0	1 deal = \$150K	2 deals = \$600K
Total ARR	\$132K	\$770K	\$2.18M

Optimistic Scenario

Strong community adoption; cloud marketplace listing in Year 2; one tier-1 healthcare or financial services OEM deal; conference exposure drives inbound enterprise pipeline.

Tier	Year 1	Year 2	Year 3
Startup (\$1K/mo)	15 customers = \$180K	40 customers = \$480K	80 customers = \$960K
Commercial (\$5K/mo)	3 customers = \$180K	10 customers = \$600K	20 customers = \$1.2M
Enterprise (\$100K/yr)	1 customer = \$100K	5 customers = \$500K	15 customers = \$1.5M
OEM / Embedded	0	1 deal = \$500K	3 deals = \$2.5M
Total ARR	\$460K	\$2.08M	\$6.16M

Highest-Leverage Early Customer Archetype

A healthcare IT platform vendor (Epic, Oracle Health, Veradigm) or a mid-size health system with an active clinical AI program. These buyers have existing regulatory compliance teams that already understand the audit trail requirement, budget authority in the \$100K–\$500K range, and internal champions (CMIOs) who are actively frustrated by LLM hallucination in clinical contexts. A single design partnership with one of these organizations provides both revenue and a reference customer that unlocks the rest of the healthcare vertical.

8. Risk Analysis

8.1 Technical Risk — 3-Hop Accuracy Gap

- **Risk:** At 3-hop reasoning depth, CEREBRUM FULL achieves 38.3% H@10 on MetaQA, compared to 72.1% for GraftNet and 81.8% for PullNet. This is a meaningful gap. In use cases requiring deep multi-hop inference (complex drug interaction chains, extended supply chain risk traversal), this gap is a competitive liability.
- **Context:** GraftNet and PullNet's 3-hop performance relies on external text retrieval — they are not pure graph reasoners. In closed-data environments (classified defense networks, air-gapped hospital systems, proprietary financial graphs), text retrieval pipelines are either unavailable or prohibited. In those environments, the 3-hop comparison is irrelevant because competing systems cannot operate at all.
- **Mitigation:** Three paths: (1) CEREBRUM OPTIMIZED is in active development and has demonstrated improvement at 1-hop (97.5%); similar optimization at 3-hop is technically plausible. (2) Hybrid pipeline — CEREBRUM handles 1-2 hop reasoning while a lightweight retrieval module handles edge cases requiring deeper traversal. (3) Strategic positioning — explicitly target use cases where 1-2 hop reasoning is sufficient (the majority of clinical decision support, financial compliance, and legal contract intelligence queries).

8.2 Market Risk — LLM Providers Adding Structured Reasoning

- **Risk:** OpenAI, Google, and Anthropic are actively adding structured reasoning capabilities to their platforms (tool use, function calling, code interpreter, structured outputs). If these capabilities mature into verifiable, graph-native reasoning, CEREBRUM's differentiation narrows.
- **Assessment:** This risk is real but slower-moving than it appears. LLM-based structured reasoning remains probabilistic at the reasoning layer even when outputs are formatted as structured data. The fundamental hallucination risk does not disappear when a model is asked to reason over a graph — it is still generating tokens, not traversing edges. The EU AI Act and FDA requirements do not accept probabilistic outputs with structured formatting as equivalent to verifiable traversal paths. The bar for regulatory compliance is higher than the bar for appearing structured.
- **Mitigation:** File provisional patent applications on DSCF and CSA architectural primitives as soon as possible. Build moat in regulated verticals where switching costs

are high and where the regulatory compliance argument is not easily replicated by LLM wrappers. Establish reference customers in healthcare and finance before LLM providers can establish compliance track records.

8.3 IP Risk — Patent Defensibility

- **Risk:** Graph attention mechanisms (Graph Attention Networks, GAT) are well-established in the literature. Broad claims on attention applied to graphs are unlikely to survive patent examination. CEREBRUM must identify and claim the novel specific combinations: DSCF community detection integrated with CSA, the training-free inference architecture, and the specific beam traversal mechanism.
- **Assessment:** A prior art search (documented by the developer as exhaustive) confirms no existing system combines training-free operation with community-structured attention for KG reasoning. The specific claim set is defensible. The risk is moderate — not that the IP is weak, but that patent prosecution takes 2–4 years and does not provide immediate protection.
- **Mitigation:** (1) File provisional patent applications immediately — provisionals establish priority date at low cost and provide 12 months of protected development time. (2) Maintain comprehensive development logs and dated documentation as trade secret protection during the prosecution window. (3) The ArXiv preprint, once published, also functions as prior art against future filers attempting to claim the same architecture.

8.4 Execution Risk — Two-Person Team

- **Risk:** Enterprise sales cycles in healthcare and finance run 6–18 months. A two-person team cannot simultaneously maintain the codebase, respond to enterprise evaluations, manage community support, and execute a conference and publication roadmap.
- **Mitigation:** Sequence activities deliberately: (1) Months 1–3: code release, ArXiv, Papers With Code — minimal enterprise sales effort. (2) Months 4–9: community adoption data as sales evidence; begin inbound enterprise conversations from research interest. (3) Month 6+: raise seed funding on community traction and one enterprise LOI; use capital to hire first enterprise account executive. This sequencing avoids the trap of spreading thin before there is enough evidence to be credible in enterprise sales rooms.

Risk Summary Matrix

Risk	Likelihood	Impact	Mitigation
3-hop accuracy gap	High (confirmed)	Medium	Optimization roadmap; hybrid pipeline; vertical targeting
LLM structured reasoning catch-up	Medium	High	Patent filing; regulated vertical moat; reference customers
IP defensibility challenge	Medium	High	Provisional patent immediately; dated dev logs; ArXiv prior art

Risk	Likelihood	Impact	Mitigation
Execution capacity (2-person team)	High	Medium	Activity sequencing; seed raise triggered by LOI; hired AE
Enterprise sales cycle length	High (structural)	Medium	AGPL flywheel pre-qualifies buyers; warm inbound preferred over cold outreach

9. Investment Thesis

Core Investment Thesis

CEREBRUM is the only commercially available AI reasoning system that is simultaneously: (a) competitive in benchmark accuracy with trained systems, (b) verifiable by architectural design, (c) deployable without training data or GPU infrastructure, and (d) natively compliant with the EU AI Act, FDA 21 CFR Part 11, SOX, HIPAA, and FINRA audit trail requirements. In an era of accelerating AI regulation, this is not a niche product. It is the baseline requirement that the mainstream market has not yet met.

9.1 Why This, Why Now

The timing is driven by three converging forces. First, the EU AI Act's transparency and audit trail requirements take full effect in 2027 — enterprise procurement cycles that begin in 2025–2026 must solve for compliance now. Second, LLM hallucination incidents in healthcare and finance have raised the organizational cost of deploying unverifiable AI to the point where procurement teams are actively looking for alternatives. Third, the knowledge graph ecosystem has matured to the point where high-quality structured data is available in healthcare (SNOMED, ICD-10, UMLS), finance (LEI, GLEIF), and biomedical research (PrimeKG, Hetionet) — the raw material CEREBRUM needs is already there.

9.2 De-Risking Milestones

- **Technical (Months 1–6):** ArXiv publication and Papers With Code submission with reproducible benchmark results. Public GitHub repository with documented installation and example notebooks. CEREBRUM OPTIMIZED 3-hop improvement results published.
- **Community (Months 3–12):** 500+ GitHub stars (early indicator of developer resonance). 50+ integrations or forks. At least one third-party reproduction of benchmark results.
- **Commercial (Months 6–18):** First paid commercial license (any tier). One signed enterprise pilot agreement (healthcare or finance preferred). Letter of Intent from a cloud marketplace or system integrator partner.

- **Fundraising Trigger:** Seed round (\$1–3M) targeting at completion of commercial milestone above. Use of funds: patent prosecution, first enterprise account executive hire, cloud marketplace listings, conference presence.

9.3 Comparable Exits and Valuations

The following transactions provide comparable reference points for a pre-revenue AI infrastructure company with novel architecture and open-source community traction:

- **Weights & Biases** — Raised \$200M Series C at \$1B valuation (2021) on developer-first AI infrastructure with OSS community flywheel. Comparable: AGPL community model, developer-led growth.
- **Hugging Face** — \$4.5B valuation (2023) on open-source ML model hub with commercial enterprise tier. Comparable: OSS flywheel converting to enterprise.
- **Neo4j** — Raised \$325M Series F at \$2B valuation (2021) on graph database with enterprise licensing. Direct comparable: graph infrastructure, dual license, regulated enterprise verticals.
- **TigerGraph** — \$1B valuation (2021) on enterprise graph analytics. Direct comparable: graph reasoning for financial and supply chain use cases.
- **Anyscale (Ray)** — Acquired by Databricks at reported \$1B+ (2024) on OSS AI infrastructure with enterprise commercial layer. Comparable: research-origin OSS to enterprise conversion.

A CEREBRUM seed round at \$5–8M pre-money valuation (consistent with the traction milestones above) with a target Series A at \$30–60M on first significant enterprise ARR (\$1–3M) is within the range established by these comparables. The path to a \$500M–\$1B exit (acquisition by a cloud provider, graph database vendor, or enterprise AI platform) is established precedent in this space.

9.4 Strategic Acquirer Landscape

Potential Acquirer	Strategic Rationale
Neo4j	Adds training-free reasoning layer to graph DB; CEREBRUM fills Neo4j's missing inference capability
Palantir	CEREBRUM's air-gap deployment and audit trail aligns directly with Palantir's government and enterprise data platform
AWS / Azure / GCP	Cloud providers competing to offer the best regulated-AI stack; CEREBRUM adds verifiable KG reasoning as a managed service
Salesforce / ServiceNow	Enterprise workflow platforms adding AI reasoning to CRM and ITSM graphs; CEREBRUM provides the reasoning layer
Oracle Health / Epic	Healthcare IT platforms needing FDA-compliant clinical reasoning; CEREBRUM is a direct clinical informatics solution

Potential Acquirer	Strategic Rationale
Databricks / Snowflake	Data lake platforms adding AI reasoning; CEREBRUM provides a structured, verifiable alternative to LLM-over-data-lake

Analyst Disclaimer & Methodology Notes

This report was prepared based on information provided by the developer of CEREBRUM (Bryan Alexander Buchorn) combined with publicly available market research, academic benchmarks, and regulatory documentation. The following assumptions and limitations apply:

- Market Size Figures:** TAM, SAM, and SOM estimates draw from Grand View Research, MarketsandMarkets, and IDC reports accessed through secondary sources. These figures are directional estimates and should be independently verified for investor presentations.
- Benchmark Comparisons:** MetaQA benchmark results for CEREBRUM are as reported by the developer. Third-party reproduction of these results has not yet been completed at the time of this report. The ArXiv submission and Papers With Code listing will provide the independent verification checkpoint.
- Revenue Projections:** All revenue projections are model-based estimates with no external validation. They represent scenario analysis, not forecasts. Actual outcomes will depend on execution, market timing, and competitive dynamics that cannot be fully anticipated at this stage.
- Regulatory Analysis:** This report does not constitute legal advice. Regulatory compliance assessments are based on publicly available regulatory text and guidance documents. Enterprise customers should obtain independent legal and compliance review before relying on CEREBRUM for regulated use cases.
- Competitive Landscape:** The competitive landscape reflects the analyst's assessment as of Q1 2026. The AI infrastructure market is moving rapidly; new entrants and product updates may alter the positioning analysis.

CEREBRUM Market Analysis Report | Bryan Alexander Buchorn (AMP, © 2026) | Prepared March 2026 | Confidential