

PARALLAX

Community-Structured Graph Attention for Knowledge Graph Reasoning

A White Paper

Authors: Bryan (Originator), Claude Sonnet 4.6 (Research Collaborator)

Date: March 2026

Status: Pre-publication draft — concept and architecture stage

License: Proprietary — all rights reserved

Abstract

ABSTRACT

We propose Parallax, a novel framework that enables Knowledge Graphs (KGs) to perform multi-hop reasoning using the same structural principles that make Transformer-based Large Language Models powerful — without requiring an LLM, without training data, and with full interpretability of every inference step. The central contribution is Community-Structured Attention (CSA): a mechanism in which graph communities serve as attention heads, graph traversal replaces matrix multiplication, and hop depth replaces layer depth. Unlike Graph Attention Networks (GATs), CSA uses community membership as a soft global constraint that captures both local topological cohesion and global structural significance simultaneously. This is made possible by a second contribution: the Dual-Signal Community Fusion (DSCF) algorithm, which produces communities that encode both LPA majority-vote structure (local) and modularity gain (global) in a single partition. We show that DSCF communities possess a structural duality that maps naturally to the dual character of multi-head attention in Transformers. Together, CSA and DSCF form an architecture where a KG can answer multi-hop questions by traversing itself, with every reasoning step grounded in explicit graph edges, every conclusion traceable to a path, and no LLM required for inference — though one may optionally be used for natural language generation.

1. Introduction

1.1 The Gap Between Knowledge Graphs and Language Models

Knowledge Graphs and Large Language Models represent two fundamentally different approaches to knowledge representation and reasoning.

Knowledge Graphs store knowledge explicitly: entities as nodes, relationships as typed edges, facts as (subject, predicate, object) triples. This makes them precise, verifiable, and updatable without retraining. However, they cannot reason beyond what is explicitly stored. Multi-hop inference — "Marie Curie discovered Polonium, Polonium is radioactive, therefore Marie Curie discovered a radioactive element" — requires either hardcoded graph traversal queries or external reasoning systems.

Large Language Models store knowledge implicitly in billions of weight parameters. They can reason, generalize, and synthesize across domains. But this implicit representation is opaque, cannot be updated without expensive fine-tuning, and is prone to hallucination — generating plausible-sounding but incorrect facts with no mechanism for ground-truth verification.

The field has responded with hybrid approaches: Retrieval-Augmented Generation (RAG), Knowledge-Graph-augmented LLMs, and GraphRAG. In all of these, the KG is a retrieval store and the LLM does the reasoning. The KG remains passive.

Parallax inverts this relationship. The KG reasons. The LLM, if present, only generates natural language from the KG's output. Every inference step is a graph traversal, every conclusion is a path, and every path can be verified.

1.2 The Core Observation

A Transformer's power comes from three mechanisms working together:

- Multi-head attention: different heads specialize on different relational aspects of the input (syntactic, semantic, long-range, etc.)
- Deep composition: each layer builds on the previous, allowing complex multi-step reasoning
- Positional awareness: the model knows where each token sits relative to others

We observe that Knowledge Graphs have natural analogs for all three:

- Community structure serves the role of attention heads: nodes within a community have strong mutual relevance; communities specialize on different conceptual domains.
- BFS hop depth serves the role of layer depth: each hop is one step of composed reasoning.
- Graph-structural features (PageRank, betweenness, degree) serve the role of positional encoding: they tell the model where each entity sits in the global information landscape.

The question is: can these analogs be made operational — not merely metaphorical? We argue yes, and demonstrate the architecture to do so.

1.3 Contributions

This paper makes three primary contributions:

- The Parallax architecture: a complete mapping of Transformer components to KG operations, enabling multi-hop reasoning via graph traversal alone.
- Community-Structured Attention (CSA): a novel attention mechanism using community membership as a soft global constraint on graph traversal, bridging the gap between local GAT-style attention and global Transformer-style attention.
- Dual-Signal Community Fusion (DSCF): a novel community detection algorithm combining LPA majority-vote and modularity gain simultaneously at each node update, producing communities with dual short-range/long-range character that maps to multi-head attention's dual specialization.

2. Background and Related Work

2.1 Graph Neural Networks

Graph Neural Networks (GNNs) [Scarselli et al., 2009] generalize neural networks to graph-structured data. The message-passing paradigm [Gilmer et al., 2017] defines node updates as aggregations of neighbor representations.

Graph Attention Networks (GATs) [Velickovic et al., 2018] introduce learned attention weights between connected nodes. Key limitations for our purposes: (a) attention is restricted to direct neighbors — no global context; (b) communities are not considered; (c) training labels required.

GraphSAGE [Hamilton et al., 2017] samples and aggregates local neighborhoods but similarly lacks global structural awareness.

None of these use community structure as an organizational principle for attention.

2.2 Knowledge Graph Reasoning

Early KG reasoning systems used rule-based approaches (AMIE [Galarraga et al., 2013]) or embedding-based methods (TransE [Bordes et al., 2013], RotatE [Sun et al., 2019]). These methods produce entity embeddings but do not perform multi-hop traversal in the attention-mechanism sense.

Path-based reasoning approaches (DeepPath [Xiong et al., 2017], MINERVA [Das et al., 2018]) use reinforcement learning to find paths. These are closer to our goal but require training data and do not use community structure.

2.3 LLM + KG Hybrid Systems

KG-GPT [Yao et al., 2023], KGPT [Chen et al., 2020], and related systems connect KGs to LLMs as retrieval stores. The LLM always performs the reasoning step; the KG is passive.

GraphRAG [Edge et al., Microsoft, 2024] uses community detection (Leiden) to summarize graph clusters as text, which is then passed to an LLM for RAG. This is the closest existing work. However:

- Communities are summarized as static text chunks, not used as attention heads
- The LLM performs all reasoning
- Paths are not returned or made interpretable
- The system is not graph-agnostic (designed for Microsoft's pipeline)

RAPTOR [Sarathi et al., 2024] builds hierarchical text clusters for RAG but operates on documents, not graphs, and uses tree structure rather than community structure.

2.4 Community Detection

The Louvain algorithm [Blondel et al., 2008] optimizes modularity greedily but can produce internally disconnected communities. Leiden [Traag et al., 2019] fixes this with a refinement phase guaranteeing internal connectivity. Label Propagation [Raghavan et al., 2007] is fast and unsupervised but non-deterministic and produces variable quality.

DSCF (this work) combines LPA and Leiden signals simultaneously at each node update, using a temperature-annealing schedule. See Section 4.2.

3. The Structural Equivalence

We establish a complete operational mapping between Transformer components and KG operations. This is not analogy — each mapping is functional.

Transformer Component	Parallax (KG) Equivalent	Notes
Token	Entity or Relation	Atomic unit of information
Vocabulary	Entity type taxonomy	Closed set of possible types
Token embedding	Entity embedding (TransE/RotatE)	Dense vector per entity
Positional encoding	Structural encoding (PR, BW, deg)	Where entity sits globally
Attention head	Community cluster (DSCF)	Specialized relational context
Attention weight	CSA weight formula	Sim + community + edge + distance
Context window	Ego-network radius R	How far to traverse
Layer depth L	BFS hop count	Reasoning step count
Feed-forward sublayer	Entity-type projection	Type-specific transformation
Residual connection	Previous-hop embedding	Prevents information loss
Layer normalization	Embedding normalization	Prevents value explosion
Output projection	Path decoder / ranker	Maps traversal to answer
KV cache	Materialized path store	Reuse traversal across queries

This equivalence has a critical implication: the number of attention heads is not a hyperparameter in Parallax — it is determined by the graph's own community structure. A graph with 12 natural communities has 12 attention heads. A graph with 200 communities has 200. The architecture adapts to the data.

4. The Parallax Architecture

4.1 Community-Structured Attention (CSA)

CSA computes attention weights for graph traversal that incorporate both local graph topology and global community structure.

Attention Weight Formula

For entity u attending to entity v at traversal hop k :

$$a(u, v, k) = \sigma($$

```

    α · cosine_sim(emb(u), emb(v))
+ β · community_score(u, v)
+ γ · edge_type_weight(type(u → v))
- δ · normalized_distance(u, v)
+ ε · hop_decay(k)
)

```

Where:

- `emb(·)` is the entity embedding (any KGE method or sentence encoder)
- `community_score(u, v)`: 1.0 if same community; 0.5 if communities are adjacent; $\exp(-\lambda \cdot \text{community_distance}(u, v))$ otherwise
- `edge_type_weight(type)`: learned or manually assigned per relation type
- `normalized_distance(u, v)`: shortest path length / graph diameter
- `hop_decay(k)`: encourages shorter paths — e.g., $1 / (1 + k)$
- σ : sigmoid activation
- $\alpha, \beta, \gamma, \delta, \epsilon$: tunable parameters

Default Parameter Values (Zero-Shot Deployment)

```

α = 0.4   (embedding similarity)
β = 0.4   (community membership)
γ = 0.1   (edge type)
δ = 0.05  (distance penalty)
ε = 0.05  (hop decay)

```

Parameters can be learned from (query, answer) pairs for supervised settings.

Why This Is Not a GAT

GATs compute $a(u, v) = f(W_u \cdot \text{emb}(u), W_v \cdot \text{emb}(v))$ — purely from learned weights on adjacent node pairs. They cannot express the community membership term $\beta \cdot \text{community_score}(u, v)$, which introduces global structural awareness without requiring the full $O(n^2)$ attention of Transformers.

CSA is $O(n \cdot \bar{k} \cdot C)$ where $\bar{k}$ is average degree and C is the average number of community-adjacent entities to consider — far cheaper than Transformer attention while capturing global structure via community membership.

4.2 Dual-Signal Community Fusion (DSCF)

DSCF is the community detection algorithm that produces the attention head structure. It is a key contribution in its own right.

The core innovation: at each individual node update, both the LPA majority-vote signal (local topology) and the modularity gain signal (global structure) are computed. The decision incorporates both simultaneously, governed by a temperature parameter:

Per-Node Update Decision

```

For each node v at each iteration:

```

```

1. LPA signal:
   lpa_cid = argmax over neighbor labels (majority vote)
   lpa_conf = vote_count[lpa_cid] / total_neighbors ∈ [0, 1]

2. Modularity signal:
   For each candidate community C adjacent to v:
        $\Delta Q(v \rightarrow C) = k_{\{v,C\}}/m - \text{resolution} \times k_v \times \sum k_C / (2m^2)$ 
   best_mod_cid = argmax  $\Delta Q$ 
   mod_conf = min(best_ΔQ × m, 1.0) ∈ [0, 1]

3. Decision:
   if lpa_cid == best_mod_cid ≠ current:
       MOVE (consensus anchor — high confidence)
   elif both say STAY:
       STAY
   elif only LPA says MOVE:
       MOVE with probability lpa_conf × temperature
   elif only modularity says MOVE:
       MOVE with probability mod_conf × (1 + (1 - temperature))
   else (disagree on different targets):
       lpa_weight = lpa_conf × temperature
       mod_weight = mod_conf × (2 - temperature)
       MOVE to weighted-random choice

temperature schedule:  $\tau_{t+1} = \max(\tau_t \times \text{cooling}, 0.01)$ 

```

Post-convergence: Leiden-style connectivity check — split any community whose induced subgraph is disconnected.

Why DSCF Communities Are the Right Attention Heads

In a trained Transformer, some heads specialize on local structure (syntactic patterns, adjacent tokens) and others on long-range structure (coreference, semantic themes). The combination allows both local and global reasoning simultaneously.

DSCF communities exhibit exactly this dual character. The LPA component ensures communities are locally coherent (nodes are topologically close). The modularity component ensures communities are globally significant (they represent structurally distinct regions of the graph). A node that is in a DSCF community is there because both local and global signals agreed.

This dual property has not previously been used as a basis for attention mechanisms in any published graph learning system.

Comparison to Leiden-Only Communities

Leiden optimizes purely for modularity. On sparse regions of a graph, Leiden may split locally coherent neighborhoods across multiple communities because the global modularity gain favors a different partition. DSCF resists this — the LPA component holds locally coherent groups together even when modularity would split them.

Comparison to LPA-Only Communities

LPA can merge structurally distinct regions if they happen to be locally connected (the "resolution limit" problem). DSCF resists this — the modularity signal penalizes over-merging.

5. The Forward Pass: Graph Reasoning

5.1 Algorithm

```

INPUT: Query Q (text string or entity list)
       Graph G (any backend)
       Community assignments C (from DSCF)
       Entity embeddings E (from any KGE method)
       Hop depth L, beam width B, top-K

OUTPUT: Ranked list of reasoning paths P = [(path, score, explanation)]

STEP 1 – Entity Grounding
  If Q is text: extract entities via NER or fuzzy match to graph vocabulary
  S = {e1, e2, ..., en}      // seed entities
  h0i = E[ei]                // initial embeddings

STEP 2 – Structural Encoding
  For each seed entity ei:
    pos_i = [pagerank(ei), betweenness(ei), degree(ei), community_id(ei)]
    h0i = LayerNorm(h0i + W_pos · pos_i)

STEP 3 – Attention Traversal (L layers)
  beam = [(path=[ei], embedding=h0i, score=1.0) for each ei in S]

  For k = 1 to L:
    candidates = []
    For each (path, h, score) in beam:
      current = path[-1]
      neighbors = G.neighbors(current)
      For each neighbor v in neighbors:
        w = CSA(current, v, k) // attention weight
        h_new = ReLU(W_k · (w · E[v] + h)) // aggregation
        h_new += h // residual
        h_new = LayerNorm(h_new)
        path_score = score × w × community_coherence(path + [v])
        candidates.append((path + [v], h_new, path_score))
    beam = top_B(candidates, key=path_score) // beam pruning

STEP 4 – Path Scoring
  Final score for path P = (e1 → r1 → e2 → r2 → ... → el):
  score(P) = Π attention_weights
            × community_coherence(P)
            × semantic_alignment(h_final, query_embedding)

STEP 5 – Output
  Return top-K paths ranked by score
  Each path includes:
  - Ordered entity/relation sequence
  - Score breakdown (attention, community, semantic)
  - Community sequence (which "heads" were traversed)
  - Natural language explanation template

```

5.2 Community Coherence

The `community_coherence` term rewards paths that traverse communities in a principled way:

```
community_coherence(path) =  
    (intra-community steps × 1.0 + cross-community steps × 0.5)  
    / total steps
```

A path that stays within one community scores 1.0 (tight local reasoning). A path that makes one community transition scores ~0.75 (one conceptual leap). This prevents paths that jump incoherently across unrelated domains.

5.3 Interpretability

The output of Parallax is always a path through the KG. This is fundamentally different from LLM reasoning in two ways:

- **Verifiable:** every step is an explicit (entity, relation, entity) triple that exists in the graph. The system cannot hallucinate a connection that is not there.
- **Auditable:** the community sequence tells you which conceptual domains were traversed. A path that crosses from "Clinical Trials" to "Drug Mechanisms" to "Side Effects" is immediately understandable.

This property makes Parallax specifically valuable in high-stakes domains (medical, legal, financial) where hallucination is unacceptable.

6. Implementation Architecture

6.1 Design Principles

Framework agnostic: Parallax must work with any graph database, any embedding method, and any LLM (or no LLM). No vendor lock-in.

No training required by default: The zero-shot configuration uses fixed parameters (α , β , γ , δ , ϵ) and any entity embedding. Communities are computed unsupervised via DSCF.

Progressive enhancement: Users can improve performance by providing training pairs (for learning parameters) or domain-specific edge type weights without changing the core architecture.

Minimal dependencies: Core (networkx, numpy, leidenalg, scipy); Adapters (optional graph DB drivers); Embeddings (optional sentence-transformers or pykeen); API (optional FastAPI).

6.2 Repository Structure

```
parallax/  
├── core/  
│   ├── graph_adapter.py      # Abstract base class for graph backends  
│   ├── embedding_engine.py   # Entity embedding interface  
│   ├── community_engine.py   # Community detection (DSCF + others)  
│   ├── attention_engine.py   # Community-Structured Attention  
│   └── structural_encoder.py  # Graph positional encoding
```

```

├── reasoning/
│   ├── traversal.py           # Beam-search attention traversal
│   ├── path_scorer.py        # Multi-signal path ranking
│   └── answer_extractor.py    # Extract top-K answers
├── adapters/
│   ├── networkx_adapter.py   # In-memory (default)
│   ├── neo4j_adapter.py      # Neo4j via bolt
│   ├── rdf_adapter.py        # SPARQL endpoint
│   └── csv_adapter.py        # Bootstrap from edge-list CSV
├── llm_bridge/
│   └── context_formatter.py   # Optional: format paths as LLM context
├── api/
│   ├── server.py             # FastAPI REST server
│   └── schemas.py            # Pydantic models
├── cli/
│   └── parallax.py           # Command-line interface
├── tests/
│   ├── test_dscf.py
│   ├── test_csa.py
│   ├── test_traversal.py
│   └── fixtures/toy_graph.csv
├── benchmarks/
│   ├── webqsp_eval.py
│   ├── metaqa_eval.py
│   └── baseline_comparison.py
├── examples/
│   ├── wikidata_quickstart.py
│   ├── neo4j_quickstart.py
│   └── csv_quickstart.py
├── pyproject.toml
├── README.md
└── PAPER.md

```

6.3 The Abstract Graph Adapter

The adapter pattern is what makes Parallax truly agnostic:

```

from abc import ABC, abstractmethod
from dataclasses import dataclass
from typing import List, Optional

@dataclass
class Entity:
    id: str
    label: str
    type: str
    properties: dict

@dataclass
class Edge:
    source_id: str
    target_id: str

```

```
relation_type: str
weight: float = 1.0

class GraphAdapter(ABC):
    @abstractmethod
    def get_entity(self, entity_id: str) -> Optional[Entity]: ...

    @abstractmethod
    def get_neighbors(self, entity_id: str,
                      edge_types: List[str] = None,
                      max_neighbors: int = 50) -> List[Edge]: ...

    @abstractmethod
    def find_entities(self, query: str, top_k: int = 10) -> List[Entity]: ...

    @abstractmethod
    def to_networkx(self) -> "nx.Graph": ...
```

Any graph system that implements these four methods works with the full Parallax stack. NetworkX, Neo4j, RDF SPARQL, property graph CSV — all handled by their respective adapter.

6.4 What Comes From AURA

The following components are directly portable from AURA with minor generalization:

- `dscf_communities()` — pure Python, depends only on `networkx`
- `leiden_communities()` / `lpa_communities()` / `hybrid_communities()`
- Neo4j connection patterns and Cypher templates
- The community broadcast WebSocket pattern (for live applications)

No other AURA-specific dependencies are carried over.

7. The DSCF-as-Attention-Head Hypothesis

The central theoretical claim of Parallax is that DSCF communities are better attention heads than Leiden-only or LPA-only communities. We state this as a falsifiable hypothesis:

```
H1 (DSCF Attention Hypothesis):
  For multi-hop reasoning tasks on KGs, Parallax with DSCF attention heads
  achieves higher answer accuracy than Parallax with Leiden-only or LPA-only
  attention heads.

H2 (CSA vs GAT Hypothesis):
  CSA-guided traversal achieves higher accuracy on multi-hop questions than
  GAT-based traversal on the same graph and same entity embeddings.

H3 (Interpretability Hypothesis):
  Parallax paths receive higher human coherence ratings than equivalent
  LLM-generated reasoning chains on the same questions, because every step
  is a grounded graph edge.
```

These hypotheses are testable on standard benchmarks (WebQSP, MetaQA-3hop) and define the empirical work for Phase 2.

8. Open Research Questions

8.1 Embedding Strategy

Two options exist:

- Option A — Pre-trained structural embeddings (TransE/RotatE): trained on the graph structure itself. More precise but requires a training step. Suitable for static KGs or when training compute is available.
- Option B — On-the-fly label embeddings (sentence-transformers): encode entity labels and descriptions using a pre-trained language model. No graph-specific training needed. More agnostic. Less precise for entities with ambiguous labels.

Recommended default: Option B for zero-shot deployment; Option A when the graph has been stable and training is feasible. Parallax should support both interchangeably via the EmbeddingEngine interface.

8.2 Adaptive Community Granularity

The DSCF resolution parameter controls how many communities are formed. Too few communities produce coarse attention heads that miss structure. Too many produce noisy heads that do not generalize.

Proposed adaptive rule: target $K \approx \sqrt{N}$ communities, where N = node count. This ensures attention head count scales sensibly with graph size:

- For AURA's KG ($N \approx 5,000$): target ~ 70 communities
- For Wikidata subset ($N \approx 100,000$): target ~ 316 communities

8.3 Soft vs Hard Community Membership

DSCF produces hard assignments (each node belongs to exactly one community). Real-world entities often span multiple communities — a person can be both a scientist and a politician.

Extension: weight-based soft membership, where each node has a probability distribution over communities. The `community_score` function becomes a dot product of membership vectors. This would require modifying DSCF to track confidence scores at convergence.

8.4 Learnable Parameters

In zero-shot mode, α , β , γ , δ , ϵ are fixed. For supervised settings, they can be learned from (query, ground-truth-answer) pairs via gradient descent on a path-ranking loss. This is an optional enhancement that does not affect the core architecture.

8.5 Temporal Knowledge Graphs

Time-stamped KGs (events, evolving relationships) introduce a temporal dimension. The positional encoding would need to incorporate temporal distance alongside graph-structural distance. Left for future work.

9. Benchmark and Evaluation Plan

9.1 Datasets

Dataset	Task	Hops	Size
WebQSP	Single + multi-hop QA	1-2	4,737 questions
MetaQA-2hop	Multi-hop QA	2	118,980 questions
MetaQA-3hop	Multi-hop QA	3	114,196 questions
FB15k-237	Link prediction	-	310,116 triples
Toy graph (internal)	Unit testing	1-4	~200 nodes

9.2 Baselines

System	Type	Notes
BFS (no attention)	Graph traversal	Traversal without CSA weighting
GAT	Graph neural network	2-layer, trained
GraphRAG	LLM-based	Community summaries to GPT-4
RAG (vanilla)	LLM-based	FAISS retrieval to GPT-4
Parallax (DSCF)	Graph attention	Ours, DSCF heads
Parallax (Leiden)	Graph attention	Ablation: Leiden-only heads
Parallax (LPA)	Graph attention	Ablation: LPA-only heads

9.3 Metrics

- Hits@1, Hits@3, Hits@10: answer in top-K paths
- Mean Reciprocal Rank (MRR): ranked answer quality
- Path coherence (human eval): are the reasoning paths understandable?
- Grounding rate: what fraction of returned paths are fully grounded (all edges verified in the KG)?

10. Broader Impact and Applications

10.1 Domain Applications

Biomedical

Drug-gene-disease-pathway graphs. Multi-hop reasoning for drug repurposing ("Drug X inhibits enzyme Y which is overexpressed in disease Z"). Grounded inference is critical — no LLM should hallucinate drug interactions.

Legal

Case law citation and statutory reference networks. Multi-hop precedent tracing. Every step in a legal argument must be citable; Parallax's grounded paths match this requirement exactly.

Cybersecurity

Attack graphs, CVE dependency networks. "What path leads from this exposed service to root access?" — a life-safety question that benefits from verified, traceable reasoning chains.

Software Engineering

Code dependency and call graphs. Impact analysis: "What does changing function X affect?" traversed as a multi-hop attention path with community context (same module = high attention).

Finance

Entity relationship graphs for regulatory compliance. Traceable reasoning chains for auditors: "Why did this transaction trigger a flag?"

10.2 The LLM Bridge

Parallax is designed to augment LLMs, not replace them. The `llm_bridge` module formats traversal output as structured context:

```
You are reasoning about: [query]

The knowledge graph traversal found these paths:

Path 1 (score: 0.94):
  Marie Curie [COMMUNITY: Scientific Discoveries]
  → [discovered] →
  Polonium [COMMUNITY: Scientific Discoveries]
  → [exhibits] →
  Radioactivity [COMMUNITY: Physics Phenomena]

Please summarize what this tells us about [query] in natural language.
```

This gives any LLM a grounded, structured context that minimizes the risk of hallucination because the facts are provided explicitly. The LLM's role is purely natural language generation, not reasoning.

10.3 The Agnosticism Property

Parallax is agnostic across five dimensions:

- Graph database: implement GraphAdapter for any system
- Embedding method: implement EmbeddingEngine for any model
- LLM: any model or none — Parallax works without one
- Domain: the algorithm is domain-blind; community structure emerges from the graph's own topology
- Query language: entities can be identified from text, IDs, or direct lookup

A single Parallax deployment can serve multiple graph backends simultaneously, which is not possible with GraphRAG or KG-specific systems.

11. Conclusion

We have presented Parallax: a framework that enables Knowledge Graphs to reason using the structural principles of Transformer attention without training data, without an LLM, and with full interpretability.

The two core contributions — Community-Structured Attention (CSA) and Dual-Signal Community Fusion (DSCF) — work together to give a KG the dual character of multi-head attention: local cohesion (from DSCF's LPA component) combined with global structural significance (from DSCF's modularity component).

The resulting system produces reasoning paths, not black-box embeddings. Every answer is traceable to a sequence of verified graph edges. Every reasoning step names the community it traversed. This interpretability property, combined with the zero-hallucination guarantee of graph-grounded inference, positions Parallax as a meaningful complement to — and in certain domains, replacement for — LLM-based reasoning over structured knowledge.

The open questions identified in Section 8 define the research program. The benchmarks in Section 9 define the empirical standard. The architecture in Section 6 defines what to build.

The name Parallax refers to the optical phenomenon where two viewpoints on the same object yield depth perception that neither viewpoint alone provides. LPA and modularity are two viewpoints on the same graph. Their combination yields structural depth — attention heads with both short-range and long-range character — that neither produces alone.

That depth is what makes the KG reason.

Appendix A: DSCF Algorithm — Full Pseudocode

```

FUNCTION dscf_communities(G, resolution=1.0, max_iter=100,
                        temp_start=1.0, cooling=0.92):

    m = |E(G)|
    IF m == 0: RETURN [{v} for v in V(G)]

    nodes = list(V(G))
    degree = {v: deg(v) for v in nodes}

    // Singleton initialization
    assignment = {v: i for i, v in enumerate(nodes)}
    com_k = {i: degree[v] for i, v in enumerate(nodes)}

    temperature = temp_start

    FOR iteration = 1 to max_iter:
        changed = False
        SHUFFLE nodes

        FOR EACH v in nodes:
            neighbors = N(v)
            IF |neighbors| == 0: CONTINUE

            cur = assignment[v]; kv = degree[v]

            // LPA signal
            vote = COUNT(assignment[nb] for nb in neighbors)
            lpa_cid = argmax(vote)
            lpa_conf = vote[lpa_cid] / |neighbors|

            // Modularity signal
            candidates = {assignment[nb] for nb in neighbors} - {cur}
            best_cid = cur; best_dq = 0
            FOR cid IN candidates:
                k_vc = |{nb in neighbors : assignment[nb] == cid}|
                dq = k_vc/m - resolution * kv * com_k[cid] / (2m2)
                IF dq > best_dq: best_dq = dq; best_cid = cid
            mod_conf = min(best_dq * m, 1.0)

            // Decision
            IF lpa_cid == best_cid ≠ cur:
                new = lpa_cid
            ELIF best_cid == cur AND lpa_cid == cur:
                CONTINUE
            ELIF best_cid == cur:
                IF random() >= lpa_conf * temperature: CONTINUE
                new = lpa_cid
            ELIF lpa_cid == cur:
                IF random() >= mod_conf * (1 + (1-temperature)): CONTINUE
                new = best_cid
            ELSE:
                lpa_w = lpa_conf * temperature
                mod_w = mod_conf * (2 - temperature)
                new = weighted_choice({lpa_cid: lpa_w, best_cid: mod_w})

            // Apply move

```



```
com_k[cur] = max(com_k[cur] - kv, 0)
com_k[new] = com_k[new] + kv
assignment[v] = new; changed = True

temperature = max(temperature * cooling, 0.01)
IF NOT changed: BREAK

// Connectivity post-pass
RETURN [component for community in assignment.values()
        for component in connected_components(G.subgraph(community))]
```

Appendix B: CSA Weight Formula — Parameter Sensitivity

The default parameter values ($\alpha=0.4$, $\beta=0.4$, $\gamma=0.1$, $\delta=0.05$, $\epsilon=0.05$) were chosen based on the following intuitions:

- Embedding similarity (α) and community membership (β) are given equal weight because both capture complementary aspects of relevance: similarity captures semantic proximity while community captures structural proximity.
- Edge type (γ) is given lower weight because it is most useful in domain-specific settings with rich edge type vocabularies.
- Distance penalty (δ) is kept small to allow multi-hop paths without excessive pruning.
- Hop decay (ϵ) is minimal to allow deep traversal when needed.

In practice, $\alpha + \beta$ dominate the attention weights for most graphs. The other parameters serve as tie-breakers and domain-adaptation handles.

Appendix C: Relationship to Existing KG Embedding Methods

TransE [Bordes et al., 2013] represents relations as translations in embedding space: $\text{emb}(h) + \text{emb}(r) \approx \text{emb}(t)$ for (h, r, t) triples. Parallax can use TransE embeddings directly for the similarity term in CSA without modification.

RotatE [Sun et al., 2019] represents relations as rotations in complex space. More expressive for symmetric, antisymmetric, and compositional relations. Also directly usable in Parallax.

Neither TransE nor RotatE produces multi-hop reasoning paths on their own. Parallax uses their embeddings as the semantic grounding layer while the traversal logic and community structure provide the reasoning.

End of white paper. Version 0.1 — March 2026.

This document is the founding specification for the Parallax project.